

Binary (Max) Heap

1. Introduction

A Binary (Max) Heap is a [complete binary tree](#) that maintains the [Max Heap property](#).

Binary Heap is one possible data structure to model an efficient [Priority Queue](#) (PQ) Abstract Data Type (ADT). In a PQ, each element has a "priority" and an element with higher priority is served before an element with lower priority (ties are broken with standard First-In First-Out (FIFO) rule as with normal Queue). Try clicking ExtractMax() for a sample animation on extracting the max value of random Binary Heap above.

To focus the discussion scope, this visualization show a Binary **Max** Heap of integers where duplicates are allowed. See [this](#) for an easy conversion to Binary **Min** Heap.

1-1. Definitions

Complete Binary Tree: Every level in the binary tree, except possibly the last/lowest level, is completely filled, and all vertices in the last level are as far left as possible.

Binary Max Heap property: The parent of each vertex - except the root - contains value greater than the value of that vertex. This is an easier-to-verify definition than the following alternative definition: The value of a vertex - except the leaf/leaves - must be greater than the value of its one (or two) child(ren).

1-2. Priority Queue ADT

Priority Queue (PQ) Abstract Data Type (ADT) is similar to normal Queue ADT, but with these two major operations:

1. Enqueue(x): Put a new element (key) **x** into the PQ (in some order),
2. **y** = Dequeue(): Return an existing element **y** that has the highest priority (key) in the PQ and if ties, return any.

Discussion: Some PQ ADT reverts to First-In First-Out (FIFO) behavior of a normal [Queue](#) in the event there is a tie of highest priority (key) in the PQ. Does guaranteeing stability on equal highest priority (key) makes PQ ADT harder to implement?

1-3. Stability of Equal Highest Key

[This is a hidden slide]

1-4. Example

Imagine: You are an [Air Traffic Controller \(ATC\)](#) working in the control tower of an airport.

You have scheduled aircraft X/Y to land in the next 3/6 minutes, respectively. Both have enough fuel for at least the next 15 minutes and both are just 2 minutes away from your airport. You observe that your airport runway is clear at the moment.

In case you do not know, aircraft can be instructed to fly in [holding pattern](#) near the airport until the designated landing time.

1-5. For Live Lecture @ NUS Only

You have to attend the live lecture to figure out what happens next...

There will be two options presented to you and you will have to decide:

1. Raise AND wave your hand if you choose option A,
2. Raise your hand but do NOT wave it if you choose option B,

If none of the two options is reasonable for you, simply do nothing.

1-6. The Example - Continued

[This is a hidden slide]

1-7. PQ Examples

There are several potential usages of PQ ADT in real-life on top of what you have seen just now (only in live lecture).

Discussion: Can you mention a few other real-life situations where a PQ is needed?

1-8. Potential Answers

[This is a hidden slide]

1-9. Linear DS for PQ?

We are able to implement this PQ ADT using (circular) array or [Linked List](#) but we will have slow (i.e., in $O(N)$) Enqueue or Dequeue operation.

Discussion: Why?

1-10. The Answer - Part 1

[This is a hidden slide]

1-11. The Answer - Part 2

[This is a hidden slide]

2. Visualisation + Max Heap Property

Now, let's view the visualisation of a (random) Binary (Max) Heap above. You should see a complete binary tree and all vertices except the root satisfy the Max Heap property ($A[\text{parent}(i)] > A[i]$ — which is still fine even with the presence of duplicate integers).

You can Toggle the Visualization Mode between the visually more intuitive complete binary tree form or the underlying compact array based implementation of a Binary (Max) Heap.

Quiz: Based on this Binary (Max) Heap property, where will the largest integer be located?

- ☐ Can be anywhere
 - ☐ At one of the leaf
 - ☐ At the root
-

2-1. Binary Heap has $O(\log N)$ Height

Important fact to memorize at this point: If we have a Binary Heap of N elements, its height will not be taller than $O(\log N)$ since we will store it as a complete binary tree.

Simple analysis: The size N of a full (more than just a complete) binary tree of height h is always $N = 2^{(h+1)} - 1$, therefore $h = \log_2(N+1) - 1 \approx \log_2 N$.

See example above with $N = 7 = 2^{(2+1)} - 1$ or $h = \log_2(7+1) - 1 = 2$.

This fact is important in the analysis of all Binary Heap-related operations.

2-2. 1-based Compact Array

A complete binary tree can be stored efficiently as a compact array A as there is no gap between vertices of a complete binary tree/elements of a compact array. To simplify the navigation operations below, we use 1-based array. VisuAlgo displays the index of each vertex as a **red label** below each vertex. Read those indices in sorted order from 1 to N , then you will see the vertices of the complete binary tree from top to down, left to right. To help you understand this, Toggle the Visualization Mode several times.

This way, we can implement basic binary tree traversal operations with simple index manipulations (with help of [bit shift manipulation](#)):

1. $\text{parent}(i) = i \gg 1$, index i divided by 2 (integer division),
2. $\text{left}(i) = i \ll 1$, index i multiplied by 2,
3. $\text{right}(i) = (i \ll 1) + 1$, index i multiplied by 2 and added by 1.

Pro tip: Try opening two copies of VisuAlgo on two browser windows. Try to visualize the same Binary Max Heap in two different modes and compare them.

3. Binary (Max) Heap Operations

In this visualization, you can perform several Binary (Max) Heap operations:

1. **Create(A)** - $O(N \log N)$ version (N calls of **Insert(v)** below)
2. **Create(A)** - $O(N)$ version
3. **Insert(v)** in $O(\log N)$ — you are allowed to insert duplicates
4. 3 versions of **ExtractMax()**:
 1. Once, in $O(\log N)$
 2. K times, i.e., **PartialSort()**, in $O(K \log N)$, or
 3. N times, i.e., **HeapSort()**, in $O(N \log N)$
5. **UpdateKey(i, newv)** in $O(\log N)$ if i is known
6. **Delete(i)** in $O(\log N)$ if i is known

There are a few other possible Binary (Max) Heap operations, but currently we do not elaborate them for pedagogical reason in a certain NUS module.

3-1. What Are The Extra Operations?

[This is a hidden slide]

4. Insert(v)

Insert(v): Insertion of a new item **v** into a Binary Max Heap can only be done at the *last index N plus 1* to maintain the compact array = complete binary tree property. However, the Max Heap property *may* still be violated. This operation then fixes Max Heap property from the insertion point **upwards** (if necessary) and stop when there is no more Max Heap property violation. Now try clicking Insert(v) several times to insert a few random **v** to the currently displayed Binary (Max Heap).

The fix Max Heap property upwards operation has no standard name. We call it **ShiftUp** but others may call it **BubbleUp** or **IncreaseKey** operation.

4-1. Why it is Correct?

Do you understand why starting from the insertion point (index **N+1**) upwards (at most until the root) and swapping a vertex with its parent when there is a Max Heap property violation during insertion is always a correct strategy?

4-2. The Answer

[This is a hidden slide]

4-3. Time Complexity Analysis

The time complexity of this **Insert(v)** operation is $O(\log N)$.

Discussion: Do you understand the derivation?

4-4. The Answer

[This is a hidden slide]

5. ExtractMax() - Once

ExtractMax(): The reporting and then the deletion of the maximum element (the root) of a Binary Max Heap requires an existing element to replace the root, otherwise the Binary Max Heap (a single complete binary tree, or 林/Lín in Chinese/tree) becomes two disjoint subtrees (two copies of 木/mù in Chinese/wood). That element

must be the *last index* N for the same reason: To maintain the compact array = complete binary tree property.

Because we promote a leaf vertex to the root vertex of a Binary Max Heap, it will very likely violates the Max Heap property. `ExtractMax()` operation then fixes Binary Max Heap property from the root **downwards** by comparing the current value with the its child/the larger of its two children (if necessary). Now try `ExtractMax()` on the currently displayed Binary (Max) Heap.

The fix Max Heap property downwards operation has no standard name. We call it **ShiftDown** but others may call it **BubbleDown** or **Heapify** operation.

5-1. Why Compare with the Larger Child?

Why if a vertex has two children, we have to check (and possibly swap) that vertex with *the larger* of its two children during the downwards fix of Max Heap property?

Why can't we just compare with the left (or right, if exists) vertex only?

5-2. The Answer

[This is a hidden slide]

5-3. Time Complexity Analysis

The time complexity of this **ExtractMax()** operation is $O(\log N)$.

Discussion: Do you understand the derivation?

5-4. The Answer

[This is a hidden slide]

6. Binary Heap for Efficient PQ

Up to here, we have a data structure that can implement the two major operations of Priority Queue (PQ) ADT efficiently:

1. For **Enqueue(x)**, we can use **Insert(x)** in $O(\log N)$ time, and
2. For $y = \text{Dequeue}()$, we can use $y = \text{ExtractMax}()$ in $O(\log N)$ time.

However, we can do a few more operations with Binary Heap.

7. Create(A) - Two Versions

Create(A): Creates a valid Binary (Max) Heap from an input array **A** of **N** integers (comma separated) into an initially empty Binary Max Heap.

There are two variants for this operations, one that is simpler but runs in $O(N \log N)$ and a more advanced technique that runs in $O(N)$.

Pro tip: Try opening two copies of VisuAlgo on two browser windows. Execute different Create(A) versions on the worst case 'Sorted example' to see the *somewhat dramatic* differences of the two.

7-1. Create(A) - $O(N \log N)$

Create(A) - $O(N \log N)$: Simply insert (that is, by calling **Insert(v)** operation) all **N** integers of the input array into an initially empty Binary Max Heap one by one.

Analysis: This operation is clearly $O(N \log N)$ as we call $O(\log N)$ **Insert(v)** operation **N** times. Let's examine the 'Sorted example' which is one of the hard case of this operation (Now try the Hard Case - $O(N \log N)$ where we show a case where **A**=[1,2,3,4,5,6,7] -- please be patient as this example will take some time to complete). If we insert values in increasing order into an initially empty Binary Max Heap, then every insertion triggers a path from the insertion point (a new leaf) upwards to the root.

7-2. Create(A) - $O(N)$

Create(A) - $O(N)$: This faster version of **Create(A)** operation was invented by Robert W. Floyd in 1964. It takes advantage of the fact that a compact array = complete binary tree and all leaves (i.e., half of the vertices — see the next slide) are Binary Max Heap by default. This operation then fixes Binary Max Heap property (if necessary) only from the last internal vertex back to the root.

Analysis: A loose analysis gives another $O(N/2 \log N) = O(N \log N)$ complexity but it is actually just $O(2*N) = O(N)$ — details in the next few slides. Now try the Hard Case - $O(N)$ on the same input array **A**=[1,2,3,4,5,6,7] and see that on the same hard case as with the previous slide (but not the one that generates maximum number of swaps), this operation is far superior than the $O(N \log N)$ version.

7-3. Many Leaf Vertices

Simple explanation on why half of Binary (Max) Heap of N (without loss of generality, let's assume that N is even) elements are leaves are as follows:

Suppose that the last leaf is at index N , then the parent of that last leaf is at index $i = N/2$ (remember [this slide](#)). The left child of vertex $i+1$, if exists (it actually does not exist), will be $2*(i+1) = 2*(N/2+1) = N+2$, which exceeds index N (the last leaf) so index $i+1$ must also be a leaf vertex that has no child. As Binary Heap indexing is consecutive, basically indices $[i+1 = N/2+1, i+2 = N/2+2, \dots, N]$, or half of the vertices, are leaves.

7-4. Why $O(N)$? - Part 1

First, we need to recall that the height of a full binary tree of size N is $\log_2 N$.

Second, we need to realise that the cost to run `shiftDown(i)` operation is not the gross upper bound $O(\log N)$, but $O(h)$ where h is the height of the subtree rooted at i .

Third, there are $\text{ceil}(N/2^{h+1})$ vertices at height h in a full binary tree.

On the example full binary tree above with $N = 7$ and $h = 2$, there are:

$\text{ceil}(7/2^{0+1}) = 4$ vertices: $\{44, 35, 26, 17\}$ at height $h = 0$,


$\text{ceil}(7/2^{1+1}) = 2$ vertices: $\{62, 53\}$ at height $h = 1$, and

$\text{ceil}(7/2^{2+1}) = 1$ vertex: $\{71\}$ at height $h = 2$.

7-5. Why $O(N)$? - Part 2

Cost of `Create(A)`, the $O(N)$ version is thus:

$$\begin{aligned}
 \underbrace{\sum_{h=0}^{\lfloor \lg(n) \rfloor} \underbrace{\left\lceil \frac{n}{2^{h+1}} \right\rceil}_{\substack{\text{\# of} \\ \text{nodes at} \\ \text{height } h}} \underbrace{O(h)}_{\substack{\text{Cost to} \\ \text{Heapify a} \\ \text{node at} \\ \text{height } h}}}_{\substack{\text{Sum over} \\ \text{all levels}}} &= \sum_{h=0}^{\lfloor \lg(n) \rfloor} \left\lceil \frac{n}{2^{h+1}} \right\rceil c \cdot h = O\left(n \sum_{h=0}^{\lfloor \lg(n) \rfloor} \frac{h}{2^h}\right) = O\left(n \sum_{h=0}^{\infty} \frac{h}{2^h}\right) = O(2N) = O(N)
 \end{aligned}$$



$$\sum_{k=0}^{\infty} kx^k = \frac{x}{(1-x)^2} = 2$$

$x = 1/2$

PS: If the formula is too complicated, a modern student can also use [WolframAlpha](#) instead.

8. HeapSort()

HeapSort(): John William Joseph Williams invented **HeapSort()** algorithm in 1964, together with this Binary Heap data structure. **HeapSort()** operation (assuming the Binary Max Heap has been created in $O(N)$) is very easy. Simply call the $O(\log N)$ **ExtractMax()** operation N times. Now try HeapSort() on the currently displayed Binary (Max) Heap.

Simple Analysis: HeapSort() clearly runs in $O(N \log N)$ — an optimal comparison-based sorting algorithm.

Quiz: **In worst case scenario, HeapSort() is asymptotically faster than...**

- ☐ Bubble Sort
- ☐ Selection Sort
- ☐ Merge Sort
- ☐ Insertion Sort

8-1. Discussion

Although **HeapSort()** runs in $\theta(N \log N)$ time for all (best/average/worst) cases, is it really the best *comparison-based* sorting algorithm?

Discussion: How about caching performance of **HeapSort()**?

8-2. The Answer

[This is a hidden slide]

8-3. PartialSort()

We can actually just call the $O(\log N)$ **ExtractMax()** operation K times if we are only interested in the top K largest elements in the Binary (Max) Heap. Now try [tobeadded] on the currently displayed Binary (Max) Heap. This operation is called PartialSort().

Simple Analysis: PartialSort() clearly runs in $O(K \log N)$ — an output-sensitive algorithm where the time complexity depends on the output size K .

9. Extras

You have reached the end of the basic stuffs of this Binary (Max) Heap data structure and we encourage you to explore further in the **Exploration Mode**.

However, we still have a few more interesting Binary (Max) Heap challenges for you that are outlined in this section.

When you have cleared them all, we invite you to study more advanced algorithms that use Priority Queue as (one of) its underlying data structure, like [Prim's MST algorithm](#), [Dijkstra's SSSP algorithm](#), A* search algorithm (not in VisuAlgo yet), a few other greedy-based algorithms, etc.

9-1. Easy Max to Min Heap Conversion

If we only deal with numbers (including this visualization that is restricted to integers only), it is easy to convert a Binary Max Heap into a Binary Min Heap without changing anything, or vice versa.

We can re-create a Binary Heap with the negation of every integer in the original Binary Heap. If we start with a Binary Max Heap, the resulting Binary Heap is a Binary Min Heap (if we ignore the negative symbols — see the picture above), and vice versa.

9-2. UpdateKey(i, newv)

For some Priority Queue applications (e.g., [HeapDecreaseKey in Dijkstra's algorithm](#)), we may have to modify (increase or decrease) the priority of an existing value that is already inserted into a Binary (Max) Heap. If the index **i** of the value is known, we can do the following simple strategy: Simply update **A[i] = newv** and then call **both shiftUp(i) and shiftDown(i)**. Only at most one of this Max Heap property restoration operation will be successful, i.e., **shiftUp(i)/shiftDown(i)** will be triggered if **newv >/< old value of A[parent(i)]/A[larger of the two children of i]**, respectively.

Thus, **UpdateKey(i, newv)** can be done in $O(\log N)$, provided we know index **i**.

9-3. Delete(i)

For some Priority Queue applications, we may have to delete an existing value that is already inserted into a Binary (Max) Heap (and this value happens to be not the root). Again, if the index **i** of the value is known, we can do the following simple strategy: Simply update **A[i] = A[1]+1** (a larger number greater than the current root),

call **shiftUp(i)** (technically, **UpdateKey(i, A[i]+1)**). This will float index **i** to be the new root, and from there, we can easily call **ExtractMax()** once to remove it.

Thus, **Delete(i)** can be done in $O(\log N)$, provided we know index **i**.

Discussion: Now for **UpdateKey(i, newv)** and **Delete(i)**, what if we are given **oldv** instead and thus we have to search for its location in the Binary (Max) Heap? Can we do this faster than $O(N)$?

9-4. The Answer

[This is a hidden slide]

9-5. Source Code

If you are looking for an implementation of Binary (Max) Heap to actually model a Priority Queue, then there is a good news.

C++ and Java already have built-in Priority Queue implementations that very likely use this data structure. They are [C++ STL priority_queue](#) (the default is a Max Priority Queue) and [Java PriorityQueue](#) (the default is a Min Priority Queue). However, the built-in implementation may not be suitable to do some PQ extended operations efficiently (details omitted for pedagogical reason in a certain NUS module).

Python [heapq](#) exists but its performance is rather slow. OCaml doesn't have built-in Priority Queue but we can use something else that is going to be mentioned in the other modules in VisuAlgo (the reason on why the details are omitted is the same as above).

PS: Heap Sort is likely used in C++ STL algorithm [partial_sort](#).

Nevertheless, here is our implementation of [BinaryHeapDemo.cpp](#).

9-6. Online Quiz

For a few more interesting questions about this data structure, please practice on [Binary Heap](#) training module (no login is required).

However, for NUS students, you should login using your official class account, officially clear this module, and such achievement will be recorded in your user account.

9-7. Online Judge Exercises

We also have a few programming problems that somewhat requires the usage of this Binary Heap data structure: [UVa 01203 - Argus](#) and [Kattis - numbertree](#).

Try them to consolidate and improve your understanding about this data structure. You are allowed to use C++ STL `priority_queue`, Python `heapq`, or Java `PriorityQueue` if that simplifies your implementation.

9-8. Discussion

[This is a hidden slide]

9-9. Shocking Stuff

After spending one long lecture on Binary (Max) Heap, here is a jaw-dropping moment...

Binary (Max) Heap data structure is probably **not** the best data structure to implement (certain operations of) ADT Priority Queue...

Discussion: So what is the alternative data structure?

9-10. The Answer

[This is a hidden slide]